



CSCI-6515 Natural Language Processing

Project 5: QA Systems with BiDAF, BERT, and RAG

Instructor: Samir Rustamov

Students: Abbas Aliyev, Huru Algayeva

Problem Statement

Part 1: Reading Comprehension

- **BiDAF Implementation**
Extract answer spans from a passage using bidirectional attention.
- **BERT Integration**
Enhance BiDAF with contextual embeddings from BERT.
- **Training & Evaluation**
Train on SQuAD-like datasets and evaluate with EM and F1 scores.

Part 2: Open-Domain Question Answering

- **Document Retrieval**
Retrieve relevant documents using sparse (TF-IDF) and dense (SBERT) methods.
- **RAG Pipeline**
Combine retrieved context with the question; extract or generate answers using BiDAF or T5/BART.
- **End-to-End Evaluation**
Measure system performance using EM, F1, BLEU, or ROUGE depending on output type.

Dataset -SQuAD version 1

A **benchmark** reading comprehension dataset with over 100,000 **question-answer pairs** on articles. Each question is answered by **extracting a span of text** from the given context. Used to train and evaluate **extractive QA models**.

| | | | |
|--------------------------|---|--|------------|
| University_of_Notre_Dame | The Joan B. Kroc Institute for International Peace Studies at the University of Notre Dame is dedica... | In what year was the Joan B. Kroc Institute for International Peace Studies founded? | 1986 |
| University_of_Notre_Dame | The Joan B. Kroc Institute for International Peace Studies at the University of Notre Dame is dedica... | To whom was John B. Kroc married? | Ray Kroc |
| University_of_Notre_Dame | The Joan B. Kroc Institute for International Peace Studies at the University of Notre Dame is dedica... | What company did Ray Kroc own? | McDonald's |
| University_of_Notre_Dame | The library system of the university is divided between the main library and each of the colleges an | How many stories tall is the main library at Notre Dame? | 14 |

Part 1 - Tokenization features

Key Fields

- **id, title** – metadata (string)
- **context, question** – raw text input
- **answers** – contains:
 - **text**: correct answer(s)
 - **answer_start**: char index in the context

Tokenized Fields

- **tokenized_context, tokenized_question** – list of tokens after tokenization
- **context_token_char_offsets** – maps each token to its original char span
- **token_answer_start, token_answer_end** – exact start/end token indices of the answer span

```
1 squad_train_tokenized.features
```

```
{'id': Value(dtype='string', id=None),
'title': Value(dtype='string', id=None),
'context': Value(dtype='string', id=None),
'question': Value(dtype='string', id=None),
'answers': Sequence(feature={'text': Value(dtype='string', id=None), 'answer_start': Value(dtype='int32', id=None)}, length=-1, id=None),
'tokenized_context': Sequence(feature=Value(dtype='string', id=None), length=-1, id=None),
'tokenized_question': Sequence(feature=Value(dtype='string', id=None), length=-1, id=None),
'context_token_char_offsets': Sequence(feature=Value(dtype='int64', id=None), length=-1, id=None),
'token_answer_start': Value(dtype='int64', id=None),
'token_answer_end': Value(dtype='int64', id=None)}
```

Vocabulary

This part creates a custom vocabulary from a tokenized SQuAD dataset:

1. Counts word frequencies in `tokenized_context` and `tokenized_question`.
2. Creates mappings:
 - o `word_to_idx` → token to ID
 - o `idx_to_word` → ID to token
 - o Includes special tokens like `<PAD>` and `<UNK>`

```
Building vocabulary...
Found 103961 unique words in the training data.
Vocabulary size (including special tokens): 103963
First 10 words in vocab: {'<PAD>': 0, '<UNK>': 1, 'the': 2, ',': 3, 'of': 4, '.': 5, 'and': 6, 'in': 7, 'to': 8, 'a': 9}
Example: ID for '<UNK>' is 1
Example: ID for 'the' is 2
Saving vocabulary to ./squad_vocab_spacy.json...
Vocabulary saved.
```

Vocabulary

Load GloVe Vectors

Reads pre-trained word embeddings (e.g., `glove.840B.300d.txt`) into a dictionary.

Build Embedding Matrix

For each word in our vocab:

- If it's in GloVe → insert the GloVe vector.
- If not → keep random vector.
- For `<PAD>` → insert all-zero vector.
- Optionally customize `<UNK>`

```
>Loading GloVe vectors from ./glove.840B.300d.txt...
Loaded 2195884 word vectors from GloVe.
Creating embedding matrix...
Embedding matrix shape: (103963, 300)
Words found in GloVe (hits): 86556
Words not found in GloVe (misses): 17407 (these will use random initialization or UNK vector)
Vector for UNK token (index 1): [-0.18693003  0.59416074 -0.40828347  0.6995341  -0.15922594  0.29638168
   0.30654564 -0.00461517 -0.2958318   0.5320532 ]...
Vector for PAD token (index 0): [0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]...
Embedding matrix converted to PyTorch tensor with shape: torch.Size([103963, 300])
Saving embedding matrix tensor to ./squad_embedding_matrix_spacy.pt...
Embedding matrix tensor saved.
```

BiDAF (Bidirectional Attention Flow) pipeline

This sets up and loads everything needed to **train the BiDAF model**:

- Loads vocab files, GloVe embeddings, and processed SQuAD datasets.
- Prepares **word + character-level inputs**.
- Defines **DataLoaders** for training and validation.
- Gets model inputs ready for batch-based training using PyTorch.

```
class CharEmbedding(nn.Module):  
    def forward(self, x_char_ids):  
        ...  
        # char CNN + max-pool  
        return final_char_emb
```

```
class BiDAFAttention(nn.Module):  
    def forward(self, C_contextual, Q_contextual, C_mask, Q_mask):  
        ...  
        return G
```

BiDAF Initialization

```
model = BiDAF(  
    word_vocab_size=WORD_VOCAB_SIZE,  
    word_embedding_dim=WORD_EMBEDDING_DIM,  
    pretrained_word_embeddings=pretrained_word_embeddings,  
    word_padding_idx=WORD_PADDING_IDX,  
    char_vocab_size=CHAR_VOCAB_SIZE,  
    char_embedding_dim=CHAR_EMBEDDING_DIM,  
    char_cnn_out_channels=CHAR_CNN_OUT_CHANNELS,  
    char_cnn_kernel_size=CHAR_CNN_KERNEL_SIZE,  
    char_padding_idx=CHAR_PADDING_IDX,  
    hidden_size=HIDDEN_SIZE,  
    num_highway_layers=NUM_HIGHWAY_LAYERS,  
    dropout_rate=DROPOUT_RATE  
).to(device)
```

```
# --- Loss Function and Optimizer ---  
criterion = nn.CrossEntropyLoss(ignore_index=-1) # Use -1 if your targets might have it for unanswerable, though SQuAD 1.1 shouldn't  
optimizer = optim.Adam(model.parameters(), lr=LEARNING_RATE)  
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', factor=0.5, patience=2, verbose=True)
```

Model Setup

Instantiates the **BiDAF model** with word & character embeddings, highway layers, LSTMs, and attention.

Loss & Optimizer

Uses [CrossEntropyLoss](#) for start/end span prediction and [Adam](#) optimizer with learning rate scheduler.

Training Loop

- Runs BiDAF forward pass on each batch
- Calculates loss for predicted start/end positions
- Applies gradient clipping to stabilize training
- Tracks progress with [tqdm](#)

Evaluation

Evaluation Loop

- Runs inference without gradients
- Computes **token-level EM and F1** for each prediction
- Saves model checkpoint with best validation F1

Train Loss: 1.2436

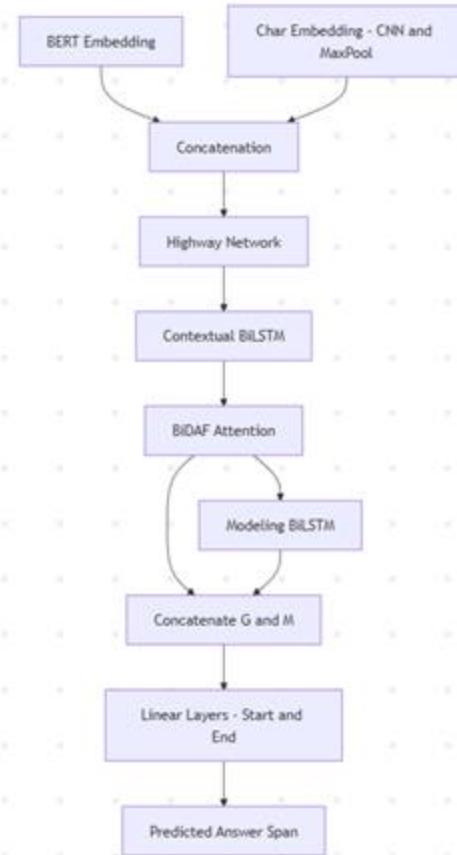
Val Loss: 4.0594 | Val EM: 42.86% | Val F1: 59.38%

Training complete.

Best Validation F1: 60.65% (Model saved at ./bidaf_best_model.pt)

BiDAF with BERT embeddings

- **CharEmbedding**
CNN + max-pooling over character tokens → captures subword patterns.
- **HighwayNetwork**
Allows the model to flexibly combine raw and transformed embeddings.
- **BiDAFAttention**
Computes attention between context and question using similarity matrix.
- **BiDAF_BERT_Char (Main Model)**
 - Inputs: BERT token IDs, attention masks, and char-level IDs.
 - BERT provides semantic context, char-CNN adds fine-grained features.
 - Outputs: start and end logits over context tokens.



Evaluation

Epoch 9 results:

exact match – 54.15%

F1 score – 72.68%

Already better than **bidaf with glove**.

```
def evaluate(model, dataloader, criterion, device):
    model.eval()
    epoch_loss = 0
    all_pred_starts, all_pred_ends = [], []
    all_true_starts, all_true_ends = [], []

    progress_bar = tqdm(dataloader, desc="Evaluating", leave=False, dynamic_ncols=True)

    with torch.no_grad():
        for batch in progress_bar:
            context_ids = batch['context_input_ids'].to(device)
            context_mask = batch['context_attention_mask'].to(device)
            context_char_ids = batch['context_token_char_ids'].to(device)
            question_ids = batch['question_input_ids'].to(device)
            question_mask = batch['question_attention_mask'].to(device)
            question_char_ids = batch['question_token_char_ids'].to(device)
            true_start_indices = batch['start_token_bert'].to(device)
            true_end_indices = batch['end_token_bert'].to(device)

            start_logits, end_logits = model(
                context_bert_ids=context_ids, context_bert_mask=context_mask, context_bert_char_ids=context_char_ids,
                question_bert_ids=question_ids, question_bert_mask=question_mask, question_bert_char_ids=question_char_ids
            )

            loss_start = criterion(start_logits, true_start_indices)
            loss_end = criterion(end_logits, true_end_indices)
            total_loss = loss_start + loss_end

            if not torch.isnan(total_loss): # Only accumulate if loss is valid
                epoch_loss += total_loss.item()

            pred_start_batch = torch.argmax(start_logits, dim=1)
            pred_end_batch = torch.argmax(end_logits, dim=1)

            all_pred_starts.extend(pred_start_batch.cpu().tolist())
            all_pred_ends.extend(pred_end_batch.cpu().tolist())
            all_true_starts.extend(true_start_indices.cpu().tolist())
            all_true_ends.extend(true_end_indices.cpu().tolist())

            progress_bar.set_postfix({'loss': f'{total_loss.item(): .4f}' if not torch.isnan(total_loss) else "NaN"})

    avg_loss = epoch_loss / len(dataloader) if len(dataloader) > 0 else 0.0
    em, f1 = compute_metrics(all_pred_starts, all_pred_ends, all_true_starts, all_true_ends)
```

Part 2 - Data extraction

This code snippet loads and preprocesses the first 30,000 examples from the SQuAD training set:

- `load_dataset("squad", split="train[:30000]")`: loads the first 30k examples from the SQuAD train split.
- `corpus`: extracts the context passages.
- `questions`: extracts the questions.
- `references`: extracts the first reference answer if available, otherwise an empty string (to handle missing answers).

```
squad = load_dataset("squad", split="train[:30000]")

# Extract all relevant fields for later evaluation
corpus = [item["context"] for item in squad]
questions = [item["question"] for item in squad]
references = [item["answers"]["text"][0] if item["answers"]["text"] else "" for item in squad] # avoid empty answers
```

Retrievers

```
# TF-IDF Sparse Retriever
class SparseRetriever:
    def __init__(self, docs):
        self.vectorizer = TfidfVectorizer().fit(docs)
        self.doc_vectors = self.vectorizer.transform(docs)
        self.docs = docs

    def retrieve(self, query, k=5):
        q_vec = self.vectorizer.transform([query])
        scores = np.dot(self.doc_vectors, q_vec.T).toarray().squeeze()
        top_k_idx = scores.argsort()[-k:][::-1]
        return [self.docs[i] for i in top_k_idx]
```

TF-IDF Sparse Retriever

- Uses `TfidfVectorizer` to convert documents and query into sparse vectors.
- Calculates cosine similarity via dot product.
- Returns top-k documents based on TF-IDF relevance.

SBERT Dense Retriever

- Uses `SentenceTransformer (multi-qa-MiniLM-L6-cos-v1)` for dense semantic embeddings.
- Applies `semantic_search` to find top-k semantically closest documents.
- More powerful for capturing meaning beyond keywords.

```
# SBERT Dense Retriever
class DenseRetriever:
    def __init__(self, docs):
        self.model = SentenceTransformer('multi-qa-MiniLM-L6-cos-v1')
        self.docs = docs
        self.doc_embeddings = self.model.encode(docs, convert_to_tensor=True)

    def retrieve(self, query, k=5):
        query_embedding = self.model.encode(query, convert_to_tensor=True)
        hits = util.semantic_search(query_embedding, self.doc_embeddings, top_k=k)[0]
        return [self.docs[hit['corpus_id']] for hit in hits]
```

T5Generator (Generative QA)

T5-base, a pre-trained encoder-decoder model, to generate answers.

- **Tokenizer**: Prepares the input (question + docs) in T5 format.
- **Model**: Generates an answer in text form.
- **Output**: Decoded text answer.

Why T5?

T5 is strong at **generative** tasks. Instead of extracting spans, it **writes answers** based on understanding context—ideal when answers aren't exact spans or when summarization helps.

```
1 # Combine documents into context
2 def augment_context(query, docs):
3     return query + " " + ".join(docs)
4
5 # Generative QA (T5 example)
6 class T5Generator:
7     def __init__(self):
8         self.tokenizer = T5Tokenizer.from_pretrained("t5-base")
9         self.model = T5ForConditionalGeneration.from_pretrained("t5-base")
10
11     def generate(self, input_text, max_len=64):
12         inputs = self.tokenizer("question: " + input_text, return_tensors="pt", truncation=True)
13         output = self.model.generate(**inputs, max_length=max_len)
14         return self.tokenizer.decode(output[0], skip_special_tokens=True)
```

BERT for Extractive QA

```
1 from transformers import AutoTokenizer, AutoModelForQuestionAnswering
2
3 # Better extractive QA model
4 tokenizer_bert = AutoTokenizer.from_pretrained("bert-large-uncased-whole-word-masking-finetuned-squad")
5 model_bert = AutoModelForQuestionAnswering.from_pretrained("bert-large-uncased-whole-word-masking-finetuned-squad")
6 model_bert.eval()
```

This snippet loads a **pretrained extractive QA model**:

- **Model:** `bert-large-uncased-whole-word-masking-finetuned-squad`
 - Fine-tuned specifically on the SQuAD dataset.
 - Outputs **start and end positions** of the answer span.
- **Why use it?**
 - It's a **stronger baseline** than vanilla BiDAF.
 - Leverages deep contextual understanding without building BiDAF from scratch.

Results

```
1 query = "Who was the first president of the United States?"  
2 answer = rag_pipeline(query, retriever, answer_mode='generate')  
3 print("Answer:", answer)
```

Answer: George Washington

```
1 query = "What is the chemical symbol for hydrogen?"  
2 answer = rag_pipeline(query, retriever, answer_mode='generate')  
3 print("Answer:", answer)
```

Answer: H

Results

```
1 query = "Where is the Eiffel Tower located?"  
2 answer = rag_pipeline(query, retriever, answer_mode='generate')  
3 print("Answer:", answer)
```

Answer: Notre Dame cathedral



```
1 query = "When did World War II end?"  
2 answer = rag_pipeline(query, retriever, answer_mode='generate')  
3 print("Answer:", answer)
```



Answer: September 11, 1776

Evaluation

SQuAD Evaluation:

Exact Match (EM): 48.00

F1 Score: 54.13

- **Exact Match (EM): 48.00**
48% of predicted answers exactly matched the ground truth span.
- **F1 Score: 54.13**
Measures overlap between predicted and actual answers (word-level). Higher than EM, as partial matches still count.

Sparse Retriever Evaluation:

Retriever Evaluation: Recall@5: 62.71%

Dense Retriever Evaluation:

Retriever Evaluation: Recall@5: 70.41%

- **Sparse Retriever (TF-IDF)**
Recall@5: 62.71%
→ Relevant doc was in the top 5 ~63% of the time.
- **Dense Retriever (SBERT)**
Recall@5: 70.41%
→ SBERT outperformed TF-IDF, finding relevant docs more reliably.

Thank you!